

19.03.2019

Statistical Methods in AI (CSE/ECE 471)

Lecture-18: Kernel Density Estimation

Ravi Kiran

Center for Visual Information Technology (CVIT), IIIT Hyderabad



ML Tasks

```
graph TD; A[ML Tasks] --> B[Predictive]; A --> C[Descriptive];
```

Predictive

Descriptive

ML::Tasks → Descriptive

- Study/Exploit the ‘structure’ of data
 - Density Estimation
 - Clustering
 - Dimensionality Reduction
- Also studied as ‘Unsupervised Learning’
 - ‘Input’ data without paired ‘Output’

Unsupervised Learning → Density Estimation

Aka “learning without a teacher”

Feature Space \mathcal{X}



Word distribution
(Probability of a word)

Task: Given $X \in \mathcal{X}$, learn $f(X)$.

$$\{x^{(1)}, \dots, x^{(m)}\} \quad x^{(i)} \in \mathbb{R}^n$$

The k -means clustering algorithm is as follows:

1. Initialize **cluster centroids** $\mu_1, \mu_2, \dots, \mu_k \in \mathbb{R}^n$ randomly.

2. Repeat until convergence: {


For every i , set

$$c^{(i)} := \arg \min_j \|x^{(i)} - \mu_j\|^2.$$

For each j , set

$$\mu_j := \frac{\sum_{i=1}^m 1\{c^{(i)} = j\} x^{(i)}}{\sum_{i=1}^m 1\{c^{(i)} = j\}}.$$

}



GMM - Algorithm

- Initialize the means μ_k , covariances Σ_k and mixing coefficients π_k

- Iterate until convergence:

- ▶ E-step: Evaluate the responsibilities given current parameters

$$\gamma_k^{(n)} = p(z^{(n)}|\mathbf{x}) = \frac{\pi_k \mathcal{N}(\mathbf{x}^{(n)}|\mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(\mathbf{x}^{(n)}|\mu_j, \Sigma_j)}$$

- ▶ M-step: Re-estimate the parameters given current responsibilities


$$\mu_k = \frac{1}{N_k} \sum_{n=1}^N \gamma_k^{(n)} \mathbf{x}^{(n)}$$

$$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^N \gamma_k^{(n)} (\mathbf{x}^{(n)} - \mu_k)(\mathbf{x}^{(n)} - \mu_k)^T$$

$$\pi_k = \frac{N_k}{N} \quad \text{with} \quad N_k = \sum_{n=1}^N \gamma_k^{(n)}$$

- ▶ Evaluate log likelihood and check for convergence

$$\ln p(\mathbf{X}|\pi, \mu, \Sigma) = \sum_{n=1}^N \ln \left(\sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}^{(n)}|\mu_k, \Sigma_k) \right)$$




spherical

Train accuracy: 88.3
Test accuracy: 92.3



diag

Train accuracy: 93.7
Test accuracy: 89.7




tied

Train accuracy: 95.5
Test accuracy: 100.0




full

Train accuracy: 94.6
Test accuracy: 97.4




setosa
versicolor
virginica

KNN




GMM





non-parametric
density estimation



- The probability that a vector x , drawn from a distribution $p(x)$, will fall in a given region \mathfrak{R} of the sample space is

$$P = \int_{\mathfrak{R}} p(x') dx'$$

- The probability that a vector x , drawn from a distribution $p(x)$, will fall in a given region \mathfrak{R} of the sample space is

$$P = \int_{\mathfrak{R}} p(x') dx'$$

- Suppose now that N vectors $\{x^{(1)}, x^{(2)}, \dots, x^{(N)}\}$ are drawn from the distribution; the probability that k of these N vectors fall in \mathfrak{R} is given by

- The probability that a vector x , drawn from a distribution $p(x)$, will fall in a given region \mathfrak{R} of the sample space is

$$P = \int_{\mathfrak{R}} p(x') dx'$$

- Suppose now that N vectors $\{x^{(1)}, x^{(2)}, \dots, x^{(N)}\}$ are drawn from the distribution; the probability that k of these N vectors fall in \mathfrak{R} is given by the binomial distribution

$$P(k) = \binom{N}{k} P^k (1 - P)^{N-k}$$

- The probability that a vector x , drawn from a distribution $p(x)$, will fall in a given region \mathfrak{R} of the sample space is

$$P = \int_{\mathfrak{R}} p(x') dx'$$

- Suppose now that N vectors $\{x^{(1)}, x^{(2)}, \dots, x^{(N)}\}$ are drawn from the distribution; the probability that k of these N vectors fall in \mathfrak{R} is given by the binomial distribution

$$P(k) = \binom{N}{k} P^k (1 - P)^{N-k}$$

- It can be shown (from the properties of the binomial p.m.f.) that the mean and variance of the ratio k/N are

$$E \left[\frac{k}{N} \right] = P \quad \text{and} \quad \text{var} \left[\frac{k}{N} \right] = E \left[\left(\frac{k}{N} - P \right)^2 \right] = \frac{P(1-P)}{N}$$

- The probability that a vector x , drawn from a distribution $p(x)$, will fall in a given region \mathfrak{R} of the sample space is

$$P = \int_{\mathfrak{R}} p(x') dx'$$

- Suppose now that N vectors $\{x^{(1)}, x^{(2)}, \dots, x^{(N)}\}$ are drawn from the distribution; the probability that k of these N vectors fall in \mathfrak{R} is given by the binomial distribution

$$P(k) = \binom{N}{k} P^k (1 - P)^{N-k}$$

- It can be shown (from the properties of the binomial p.m.f.) that the mean and variance of the ratio k/N are

$$E \left[\frac{k}{N} \right] = P \quad \text{and} \quad \text{var} \left[\frac{k}{N} \right] = E \left[\left(\frac{k}{N} - P \right)^2 \right] = \frac{P(1-P)}{N}$$

- Therefore, as $N \rightarrow \infty$ the distribution becomes sharper (the variance gets smaller), so we can expect that a good estimate of the probability P can be obtained from the mean fraction of the points that fall within \mathfrak{R}

$$P \cong \frac{k}{N}$$

- On the other hand, if we assume that \mathfrak{R} is so small that $p(x)$ does not vary appreciably within it, then

$$\int_{\mathfrak{R}} p(x') dx' \cong p(x)V$$

- where V is the volume enclosed by region \mathfrak{R}

- On the other hand, if we assume that \mathfrak{R} is so small that $p(x)$ does not vary appreciably within it, then

$$\int_{\mathfrak{R}} p(x') dx' \cong p(x)V$$

- where V is the volume enclosed by region \mathfrak{R}

- Merging with the previous result we obtain

$$\left. \begin{array}{l} P = \int_{\mathfrak{R}} p(x') dx' \cong p(x)V \\ P \cong \frac{k}{N} \end{array} \right\} \Rightarrow p(x) \cong \frac{k}{NV}$$

The histogram


The simplest form of non-parametric DE is the histogram

- Divide the sample space into a number of bins and approximate the density at the center of each bin by the fraction of points in the training data that fall into the corresponding bin

$$p(x) \cong \frac{k}{NV}$$



$$p_H(x) = \frac{1}{N} \frac{[\# \text{ of } x^{(k)} \text{ in same bin as } x]}{[\text{width of bin}]}$$

- The histogram requires two “parameters” to be defined: bin width and starting position of the first bin




$$p_H(x) = \frac{1}{N} \frac{[\# \text{ of } x^{(k)} \text{ in same bin as } x]}{[\text{width of bin}]}$$

$$P_{dx} \approx f(x)dx = \frac{N_{dx}}{N}$$





A toy example

- (the log of) wing spans of aircraft built from 1956 – 1984
- Wing-spans: 2, 22, 42, 62, 82, 102, 122, 142, 162, 182, 202, 222



A toy example

- (the log of) wing spans of aircraft built from 1956 – 1984
- Wing-spans: 2, 22, 42, 62, 82, 102, 122, 142, 162, 182, 202, 222




Issues with Histograms

- Not smooth
- Depend on end points of bins
- Depend on width of bins

Characterizing density – a data-driven perspective

Box Kernel Density Estimate

- 12 data-points
- block width $\frac{1}{2}$, height $\frac{1}{6}$




- In conclusion, the general expression for non-parametric density estimation becomes

$$p(x) \cong \frac{k}{NV} \text{ where } \begin{cases} V & \text{volume surrounding } x \\ N & \text{total \#examples} \\ k & \text{\#examples inside } V \end{cases}$$

Parzen windows

Problem formulation


- Assume that the region \mathfrak{R} that encloses the k examples is a hypercube with sides of length h centered at x
 - Then its volume is given by $V = h^D$, where D is the number of dimensions



Parzen windows

Problem formulation

- Assume that the region \mathfrak{R} that encloses the k examples is a hypercube with sides of length h centered at x
 - Then its volume is given by $V = h^D$, where D is the number of dimensions



- To find the number of examples that fall within this region we define a kernel function $K(u)$

$$K(u) = \begin{cases} 1 & |u_j| < 1/2 \quad \forall j = 1 \dots D \\ 0 & \text{otherwise} \end{cases}$$

- This kernel, which corresponds to a unit hypercube centered at the origin, is known as a Parzen window or the naïve estimator
- The quantity $K((x - x^{(n)})/h)$ is then equal to unity if $x^{(n)}$ is inside a hypercube of side h centered on x , and zero otherwise


- The total number of points inside the hypercube is then

$$k = \sum_{n=1}^N K\left(\frac{x - x^{(n)}}{h}\right)$$

Substituting back into the expression for the density estimate

$$p_{KDE}(x) = \frac{1}{Nh^D} \sum_{n=1}^N K\left(\frac{x - x^{(n)}}{h}\right)$$


- Notice how the Parzen window estimate resembles the histogram, with the exception that the bin locations are determined by the data



Exercise

- Given dataset $X = \{4, 5, 5, 6, 12, 14, 15, 15, 16, 17\}$, use Parzen windows to estimate the density $p(x)$ at $y = 3, 10, 15$; use $h = 4$
- Solution

- Let's first draw the dataset to get an idea of the data




- Let's now estimate $p(y = 3)$

Exercise

- Given dataset $X = \{4, 5, 5, 6, 12, 14, 15, 15, 16, 17\}$, use Parzen windows to estimate the density $p(x)$ at $y = 3, 10, 15$; use $h = 4$
- Solution

- Let's first draw the dataset to get an idea of the data




- Let's now estimate $p(y = 3)$

$$p(y = 3) = \frac{1}{Nh^D} \sum_{n=1}^N K\left(\frac{x - x^{(n)}}{h}\right)$$

Exercise

- Given dataset $X = \{4, 5, 5, 6, 12, 14, 15, 15, 16, 17\}$, use Parzen windows to estimate the density $p(x)$ at $y = 3, 10, 15$; use $h = 4$
- Solution

- Let's first draw the dataset to get an idea of the data




- Let's now estimate $p(y = 3)$

$$p(y = 3) = \frac{1}{Nh^D} \sum_{n=1}^N K\left(\frac{x - x^{(n)}}{h}\right) = \frac{1}{10 \times 4^1} \left[K\left(\frac{3-4}{4}\right) + K\left(\frac{3-5}{4}\right) + \dots + K\left(\frac{3-17}{4}\right) \right] = 0.0025$$

Exercise

- Given dataset $X = \{4, 5, 5, 6, 12, 14, 15, 15, 16, 17\}$, use Parzen windows to estimate the density $p(x)$ at $y = 3, 10, 15$; use $h = 4$
- Solution

- Let's first draw the dataset to get an idea of the data



- Let's now estimate $p(y = 3)$

$$p(y = 3) = \frac{1}{Nh^D} \sum_{n=1}^N K\left(\frac{x - x^{(n)}}{h}\right) = \frac{1}{10 \times 4^1} \left[K\left(\frac{3-4}{4}\right) + K\left(\frac{3-5}{4}\right) + \dots + K\left(\frac{3-17}{4}\right) \right] = 0.0025$$

- Similarly

$$p(y = 10) = \frac{1}{10 \times 4^1} [0 + 0 + 0 + 0 + 0 + 0 + 0 + 0 + 0 + 0] = 0$$

$$p(y = 15) = \frac{1}{10 \times 4^1} [0 + 0 + 0 + 0 + 0 + 0 + 1 + 1 + 1 + 1 + 0] = 0.1$$

Smooth kernels

The Parzen window has several drawbacks

- It yields density estimates that have discontinuities
- It weights equally all points x_i , regardless of their distance to the estimation point x

Smooth kernels

The Parzen window has several drawbacks



- It yields density estimates that have discontinuities
- It weights equally all points x_i , regardless of their distance to the estimation point x

For these reasons, the Parzen window is commonly replaced with a smooth kernel function $K(u)$

$$\int_{\mathbb{R}^D} K(x) dx = 1$$

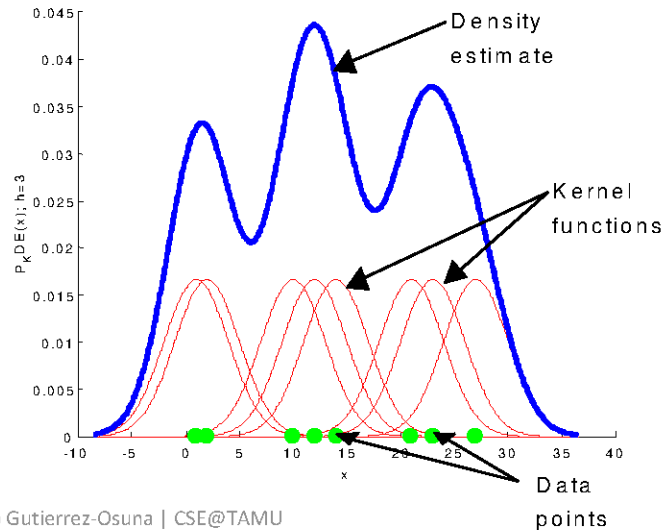
- Usually, but not always, $K(u)$ will be a radially symmetric and unimodal pdf, such as the Gaussian $K(x) = (2\pi)^{-D/2} e^{-\frac{1}{2}x^T x}$
- Which leads to the density estimate

$$p_{KDE}(x) = \frac{1}{Nh^D} \sum_{n=1}^N K\left(\frac{x - x^{(k)}}{h}\right)$$



Interpretation





- Just as the Parzen window estimate can be seen as a sum of boxes centered at the data, the smooth kernel estimate is a sum of “bumps”
- The kernel function determines the shape of the bumps
- The parameter h , also called the smoothing parameter or bandwidth, determines their width



Bandwidth selection

The problem of choosing h is crucial in density estimation

- A large h will over-smooth the DE and mask the structure of the data
- A small h will yield a DE that is spiky and very hard to interpret



- We would like to find a value of h that minimizes the error between the estimated density and the true density

- A natural measure is the MSE at the estimation point x , defined by

$$E[(p_{KDE}(x) - p(x))^2] = E[p_{KDE}(x) - p(x)]^2 + var(p_{KDE}(x))$$

- We would like to find a value of h that minimizes the error between the estimated density and the true density

- A natural measure is the MSE at the estimation point x , defined by


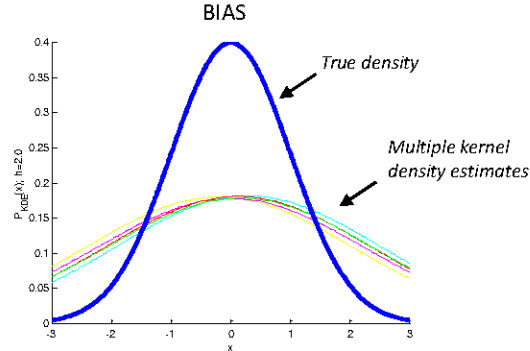
$$E[(p_{KDE}(x) - p(x))^2] = \underbrace{E[p_{KDE}(x) - p(x)]^2}_{bias} + \underbrace{var(p_{KDE}(x))}_{variance}$$

- This expression is an example of the bias-variance tradeoff that we saw in an earlier lecture: the bias can be reduced at the expense of the variance, and vice versa
 - The bias of an estimate is the systematic error incurred in the estimation
 - The variance of an estimate is the random error incurred in the estimation

- The bias-variance dilemma applied to bandwidth selection simply means that

– The bias-variance dilemma applied to bandwidth selection simply means that

- A large bandwidth will reduce the differences among the estimates of $p_{KDE}(x)$ for different data sets (the variance), but it will increase the bias of $p_{KDE}(x)$ with respect to the true density $p(x)$
- A small bandwidth will reduce the bias of $p_{KDE}(x)$, at the expense of a larger variance in the estimates $p_{KDE}(x)$



Bandwidth selection methods, univariate case

Subjective choice

- The natural way for choosing h is to plot out several curves and choose the estimate that best matches one's prior (subjective) ideas
- However, this method is not practical in pattern recognition since we typically have high-dimensional data

Reference to a standard distribution

- Assume a standard density function and find the value of the bandwidth that minimizes the integral of the square error (MISE)

$$h_{MISE} = \arg \min \{E[\int (p_{KDE}(x) - p(x))^2 dx]\}$$

- If we assume that the true distribution is Gaussian and we use a Gaussian kernel, it can be shown that the optimal value of h is

$$h^* = 1.06\sigma N^{-1/5}$$

- where σ is the sample standard deviation and N is the number of training examples

- Better results can be obtained by
 - Using a robust measure of the spread instead of the sample variance, and
 - Reducing the coefficient 1.06 to better cope with multimodal densities
 - The optimal bandwidth then becomes

$$h^* = 0.9AN^{-1/5} \text{ where } A = \min\left(\sigma, \frac{IQR}{1.34}\right)$$

- IQR is the interquartile range, a robust estimate of the spread
 - IQR is the difference between the 75th percentile ($Q3$) and the 25th percentile ($Q1$): $IQR = Q3 - Q1$
 - A percentile rank is the proportion of examples in a distribution that a specific example is greater than or equal to

Maximum likelihood cross-validation





Maximum likelihood cross-validation

- The ML estimate of h is degenerate since it yields $h_{ML} = 0$, a density estimate with Dirac delta functions at each training data point

Maximum likelihood cross-validation

- The ML estimate of h is degenerate since it yields $h_{ML} = 0$, a density estimate with Dirac delta functions at each training data point
- A practical alternative is to maximize the “pseudo-likelihood” computed using leave-one-out cross-validation

$$h^* = \arg \max \left\{ \frac{1}{N} \sum_{n=1}^N \log p_{-n}(x^{(n)}) \right\}$$
$$\text{where } p_{-n}(x^{(n)}) = \frac{1}{(N-1)h} \sum_{\substack{m=1 \\ m \neq n}}^N K\left(\frac{x^{(n)} - x^{(m)}}{h}\right)$$



Multivariate density estimation

For the multivariate case, the KDE is

$$p_{KDE}(x) = \frac{1}{Nh^D} \sum_{n=1}^N K\left(\frac{x-x^{(n)}}{h}\right)$$

- Notice that the bandwidth h is the same for all the axes, so this density estimate will be weight all the axis equally
- If one or several of the features has larger spread than the others, we should use a vector of smoothing parameters or even a full covariance matrix, which complicates the procedure

Multivariate density estimation


For the multivariate case, the KDE is

$$p_{KDE}(x) = \frac{1}{Nh^D} \sum_{n=1}^N K\left(\frac{x-x^{(n)}}{h}\right)$$

- Notice that the bandwidth h is the same for all the axes, so this density estimate will be weight all the axis equally
- If one or several of the features has larger spread than the others, we should use a vector of smoothing parameters or even a full covariance matrix, which complicates the procedure

There are two basic alternatives to solve the scaling problem without having to use a more general KDE

- Pre-scaling each axis (normalize to unit variance, for instance)
- Pre-whitening the data (linearly transform so $\Sigma = I$), estimate the density, and then transform back [Fukunaga]
 - The whitening transform is $y = \Lambda^{-1/2} M^T x$, where Λ and M are the eigenvalue and eigenvector matrices of Σ
 - Fukunaga's method is equivalent to using a hyper-ellipsoidal kernel



Product kernels

A good alternative for multivariate KDE is the product kernel

$$p_{PKDE}(x) = \frac{1}{N} \sum_{i=1}^N K(x, x^{(n)}, h_1, \dots, h_D)$$

$$\text{where } K(x, x^{(n)}, h_1, \dots, h_D) = \frac{1}{h_1 \dots h_D} \prod_{d=1}^D K_d \left(\frac{x_d - x_d^{(n)}}{h_d} \right)$$

- The product kernel consists of the product of one-dimensional kernels
 - Typically the same kernel function is used in each dimension ($K_d(x) = K(x)$), and only the bandwidths are allowed to differ
 - Bandwidth selection can then be performed with any of the methods presented for univariate density estimation

Product kernels

A good alternative for multivariate KDE is the product kernel

$$p_{PKDE}(x) = \frac{1}{N} \sum_{i=1}^N K(x, x^{(n)}, h_1, \dots, h_D)$$

$$\text{where } K(x, x^{(n)}, h_1, \dots, h_D) = \frac{1}{h_1 \dots h_D} \prod_{d=1}^D K_d \left(\frac{x_d - x_d^{(n)}}{h_d} \right)$$

- The product kernel consists of the product of one-dimensional kernels
 - Typically the same kernel function is used in each dimension ($K_d(x) = K(x)$), and only the bandwidths are allowed to differ
 - Bandwidth selection can then be performed with any of the methods presented for univariate density estimation

$K(x, x^{(n)}, h_1, \dots, h_D)$ uses kernel independence
features are independent

Product kernels

A good alternative for multivariate KDE is the product kernel

$$p_{PKDE}(x) = \frac{1}{N} \sum_{i=1}^N K(x, x^{(n)}, h_1, \dots, h_D)$$

$$\text{where } K(x, x^{(n)}, h_1, \dots, h_D) = \frac{1}{h_1 \dots h_D} \prod_{d=1}^D K_d \left(\frac{x_d - x_d^{(n)}}{h_d} \right)$$

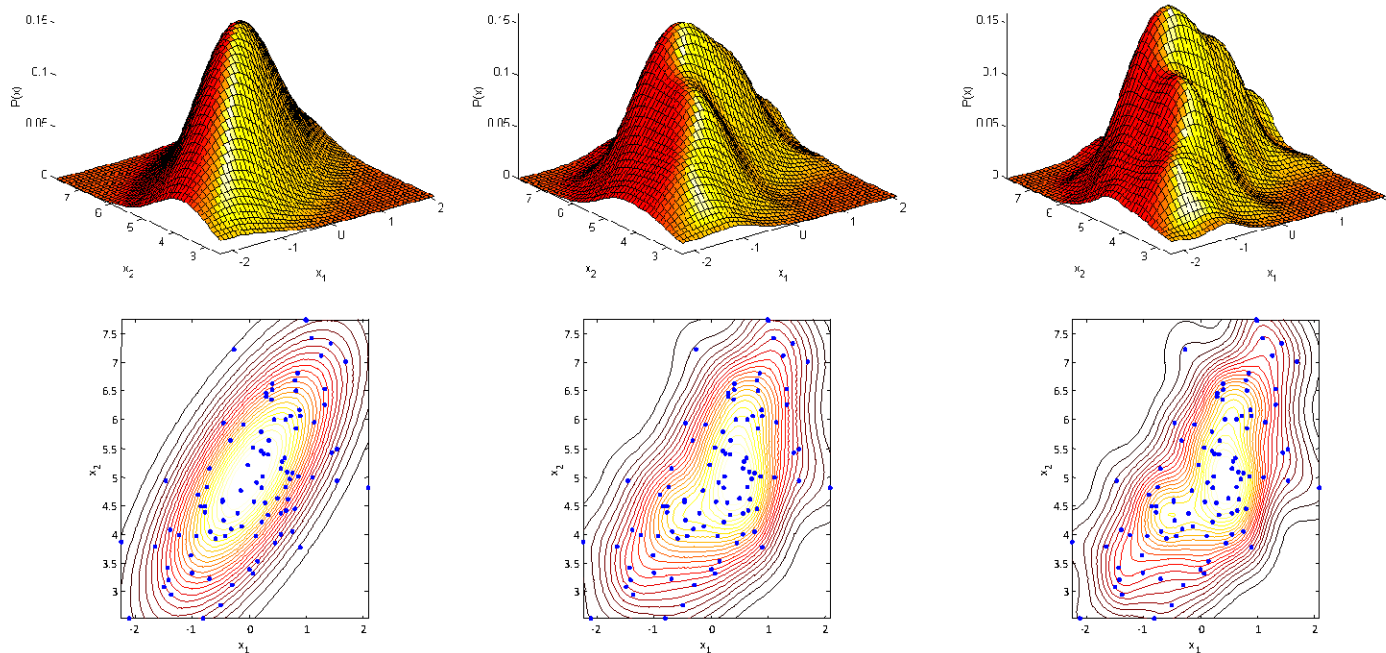
- The product kernel consists of the product of one-dimensional kernels
 - Typically the same kernel function is used in each dimension ($K_d(x) = K(x)$), and only the bandwidths are allowed to differ
 - Bandwidth selection can then be performed with any of the methods presented for univariate density estimation
- Note that although $K(x, x^{(n)}, h_1, \dots, h_D)$ uses kernel independence does not imply we assume the features are independent
 - If we assumed feature independence, the DE would have the expression

$$p_{FEAT-IND}(x) = \prod_{d=1}^D \frac{1}{N h^D} \sum_{i=1}^N K_d \left(\frac{x_d - x_d^{(n)}}{h_d} \right)$$

- Notice how the order of the summation and product are reversed compared to the product kernel

Example I

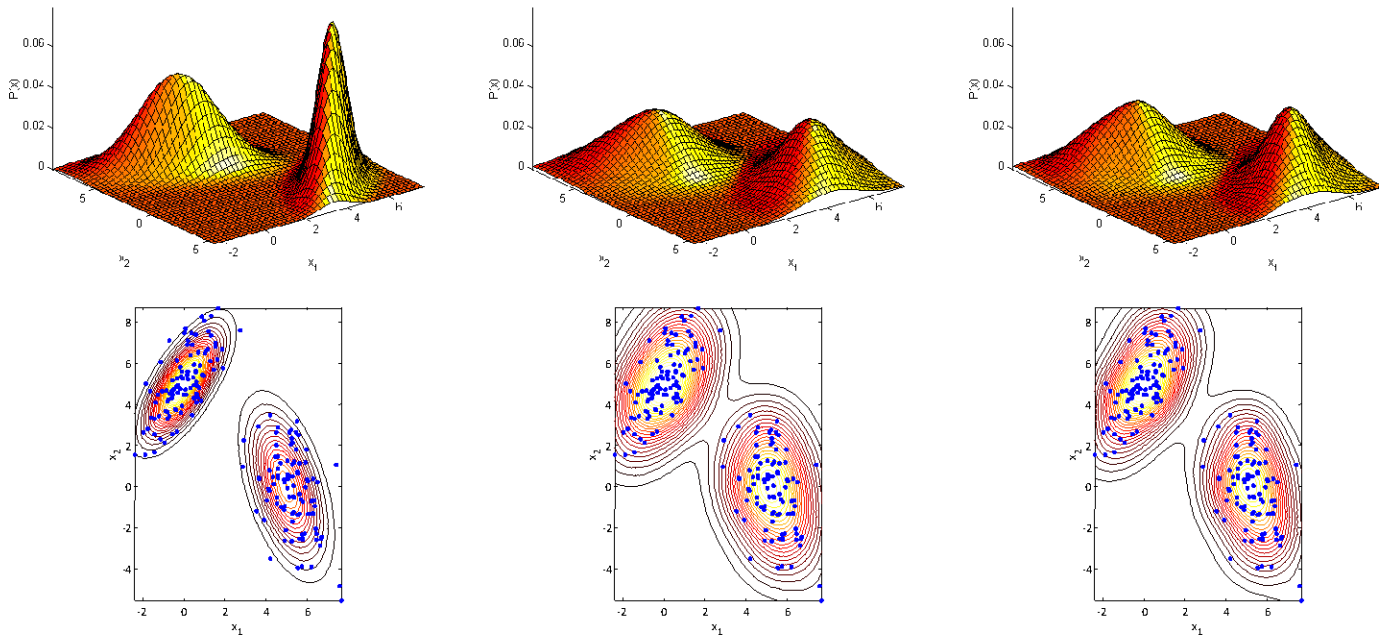
- This example shows the product KDE of a bivariate unimodal Gaussian
 - 100 data points were drawn from the distribution
 - The figures show the true density (left) and the estimates using $h = 1.06\sigma N^{-1/5}$ (middle) and $h = 0.9AN^{-1/5}$ (right)



Example II

– This example shows the product KDE of a bivariate bimodal Gaussian

- 100 data points were drawn from the distribution
- The figures show the true density (left) and the estimates using $h = 1.06\sigma N^{-1/5}$ (middle) and $h = 0.9AN^{-1/5}$ (right)



$$p(x) \cong \frac{k}{NV} \text{ where } \begin{cases} V & \text{volume surrounding } x \\ N & \text{total \#examples} \\ k & \text{\#examples inside } V \end{cases}$$

- We can fix V and determine k from the data. This leads to **kernel density estimation** (KDE), the subject of this lecture
- We can fix k and determine V from the data. This gives rise to the **k-nearest-neighbor** (kNN) approach

KDE

- <https://scikit-learn.org/stable/modules/density.html>

```
>>> from sklearn.neighbors.kde import KernelDensity
>>> import numpy as np
>>> X = np.array([[ -1, -1], [-2, -1], [-3, -2], [ 1,  1], [ 2,  1], [ 3,  2]])
>>> kde = KernelDensity(kernel='gaussian', bandwidth=0.2).fit(X)
>>> kde.score_samples(X)
array([-0.41075698, -0.41075698, -0.41076071, -0.41075698, -0.41075698,
       -0.41076071])
```

References

- <https://mathisonian.github.io/kde/>
- [https://en.wikipedia.org/wiki/Kernel density estimation](https://en.wikipedia.org/wiki/Kernel_density_estimation) (Sections 1,2)
- Sampling from distributions:
<http://karlrosaen.com/ml/notebooks/simulating-random-variables/>